

Name : Vrajkumar Patel

House Prices - Advanced Regression Techniques

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split
```

```
#Loading train and test datasets
train_data = pd.read_csv('./train.csv')
test_data = pd.read_csv('./test.csv')
```

#Viewing the contents of the dataset

train data

LotShape	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
0 Reg	1	60	RL	65.0	8450	Pave	NaN
1 Reg	2	20	RL	80.0	9600	Pave	NaN
2 IR1	3	60	RL	68.0	11250	Pave	NaN
3 IR1	4	70	RL	60.0	9550	Pave	NaN
4 IR1	5	60	RL	84.0	14260	Pave	NaN
...
1455 Reg	1456	60	RL	62.0	7917	Pave	NaN
1456 Reg	1457	20	RL	85.0	13175	Pave	NaN
1457 Reg	1458	70	RL	66.0	9042	Pave	NaN
1458 Reg	1459	20	RL	68.0	9717	Pave	NaN
1459 Reg	1460	20	RL	75.0	9937	Pave	NaN

MiscVal	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature
0	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
1	Lvl	AllPub	...	0	NaN	NaN	NaN
0							

2	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
3	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
4	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
...
...							
1455	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
1456	Lvl	AllPub	...	0	NaN	MnPrv	NaN
0							
1457	Lvl	AllPub	...	0	NaN	GdPrv	Shed
2500							
1458	Lvl	AllPub	...	0	NaN	NaN	NaN
0							
1459	Lvl	AllPub	...	0	NaN	NaN	NaN
0							

	MoSold	YrSold	SaleType	SaleCondition	SalePrice
0	2	2008	WD	Normal	208500
1	5	2007	WD	Normal	181500
2	9	2008	WD	Normal	223500
3	2	2006	WD	Abnorml	140000
4	12	2008	WD	Normal	250000
...
1455	8	2007	WD	Normal	175000
1456	2	2010	WD	Normal	210000
1457	5	2010	WD	Normal	266500
1458	4	2010	WD	Normal	142125
1459	6	2008	WD	Normal	147500

[1460 rows x 81 columns]

By getting the info on train data, we can see the columns and their datatypes. Our task is to predict housing prices using the attributes **Square footage, No. of Bedrooms & No. of Bathrooms**

```
train_data.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1460 entries, 0 to 1459
Data columns (total 81 columns):
#   Column          Non-Null Count  Dtype
---  -
0   Id              1460 non-null   int64
1   MSSubClass      1460 non-null   int64
2   MSZoning        1460 non-null   object
3   LotFrontage     1201 non-null   float64
```

4	LotArea	1460	non-null	int64
5	Street	1460	non-null	object
6	Alley	91	non-null	object
7	LotShape	1460	non-null	object
8	LandContour	1460	non-null	object
9	Utilities	1460	non-null	object
10	LotConfig	1460	non-null	object
11	LandSlope	1460	non-null	object
12	Neighborhood	1460	non-null	object
13	Condition1	1460	non-null	object
14	Condition2	1460	non-null	object
15	BldgType	1460	non-null	object
16	HouseStyle	1460	non-null	object
17	OverallQual	1460	non-null	int64
18	OverallCond	1460	non-null	int64
19	YearBuilt	1460	non-null	int64
20	YearRemodAdd	1460	non-null	int64
21	RoofStyle	1460	non-null	object
22	RoofMatl	1460	non-null	object
23	Exterior1st	1460	non-null	object
24	Exterior2nd	1460	non-null	object
25	MasVnrType	588	non-null	object
26	MasVnrArea	1452	non-null	float64
27	ExterQual	1460	non-null	object
28	ExterCond	1460	non-null	object
29	Foundation	1460	non-null	object
30	BsmtQual	1423	non-null	object
31	BsmtCond	1423	non-null	object
32	BsmtExposure	1422	non-null	object
33	BsmtFinType1	1423	non-null	object
34	BsmtFinSF1	1460	non-null	int64
35	BsmtFinType2	1422	non-null	object
36	BsmtFinSF2	1460	non-null	int64
37	BsmtUnfSF	1460	non-null	int64
38	TotalBsmtSF	1460	non-null	int64
39	Heating	1460	non-null	object
40	HeatingQC	1460	non-null	object
41	CentralAir	1460	non-null	object
42	Electrical	1459	non-null	object
43	1stFlrSF	1460	non-null	int64
44	2ndFlrSF	1460	non-null	int64
45	LowQualFinSF	1460	non-null	int64
46	GrLivArea	1460	non-null	int64
47	BsmtFullBath	1460	non-null	int64
48	BsmtHalfBath	1460	non-null	int64
49	FullBath	1460	non-null	int64
50	HalfBath	1460	non-null	int64
51	BedroomAbvGr	1460	non-null	int64
52	KitchenAbvGr	1460	non-null	int64

```

53 KitchenQual      1460 non-null object
54 TotRmsAbvGrd     1460 non-null int64
55 Functional        1460 non-null object
56 Fireplaces        1460 non-null int64
57 FireplaceQu       770 non-null object
58 GarageType        1379 non-null object
59 GarageYrBlt       1379 non-null float64
60 GarageFinish      1379 non-null object
61 GarageCars        1460 non-null int64
62 GarageArea        1460 non-null int64
63 GarageQual        1379 non-null object
64 GarageCond        1379 non-null object
65 PavedDrive        1460 non-null object
66 WoodDeckSF        1460 non-null int64
67 OpenPorchSF       1460 non-null int64
68 EnclosedPorch     1460 non-null int64
69 3SsnPorch         1460 non-null int64
70 ScreenPorch       1460 non-null int64
71 PoolArea          1460 non-null int64
72 PoolQC            7 non-null object
73 Fence            281 non-null object
74 MiscFeature       54 non-null object
75 MiscVal          1460 non-null int64
76 MoSold           1460 non-null int64
77 YrSold            1460 non-null int64
78 SaleType          1460 non-null object
79 SaleCondition     1460 non-null object
80 SalePrice         1460 non-null int64
dtypes: float64(3), int64(35), object(43)
memory usage: 924.0+ KB

```

We can see that the no. of bathrooms is provided in 2 varieties: **Full Bathrooms & Half Bathrooms** and there are 2 columns for each variant i.e. For Full Bathrooms : **FullBath & BsmtFullBath**, For Half Bathrooms : **HalfBath & BsmtHalfBath**

We can find the correlation between the various bathrooms and sales price

```

bath_sales_df = train_data[['FullBath', 'BsmtFullBath', 'HalfBath',
                             'BsmtHalfBath', 'SalePrice']]
bath_sales_df


```

	FullBath	BsmtFullBath	HalfBath	BsmtHalfBath	SalePrice
0	2	1	1	0	208500
1	2	0	0	1	181500
2	2	1	1	0	223500
3	1	1	0	0	140000
4	2	1	1	0	250000
...
1455	2	0	1	0	175000

1456	2	1	0	0	210000
1457	2	0	0	0	266500
1458	1	1	0	0	142125
1459	1	1	1	0	147500


[1460 rows x 5 columns]

```
sns.heatmap(bath_sales_df.corr(), annot=True)
plt.show()
```



```
#We can create new features for each bathroom type (Full or Half)
fullcomb = train_data['FullBath'] + train_data['BsmtFullBath']
halfcomb = train_data['HalfBath'] + train_data['BsmtHalfBath']

bath_sales_new = pd.DataFrame({'FullBathNew':fullcomb,
                               'HalfBathNew':halfcomb, 'SalePrice':train_data['SalePrice']})
sns.heatmap(bath_sales_new.corr(), annot=True)
plt.show()
```



We can see the correlation is improved with new variables compared to older ones

#We need only Square footage, No.of Bedrooms and No.of Bathrooms data for our ML model

```
train_data_filt = pd.DataFrame({'LotArea':train_data['LotArea'],
                                'FullBathComb':fullcomb, 'HalfBathComb':halfcomb,
                                'Bedrooms':train_data['BedroomAbvGr'],
                                'SalePrice':train_data['SalePrice']})
train_data_filt.head()
```

	LotArea	FullBathComb	HalfBathComb	Bedrooms	SalePrice
0	8450	3	1	3	208500
1	9600	2	1	3	181500
2	11250	3	1	3	223500
3	9550	2	0	3	140000
4	14260	3	1	4	250000


We can see that it has no null values

```
train_data_filt.isna().sum()
```


```
LotArea      0
FullBathComb 0
HalfBathComb 0
Bedrooms     0
SalePrice    0
dtype: int64
```

We can see the distribution of input parameters against sale price

```
train_data_filt.plot(kind='scatter',x='LotArea', y='SalePrice', )  
<Axes: xlabel='LotArea', ylabel='SalePrice'>
```



```
train_data_filt.plot(kind='scatter',x='Bedrooms', y='SalePrice' )  
<Axes: xlabel='Bedrooms', ylabel='SalePrice'>
```



```
train_data_filt.describe()
```

	LotArea	FullBathComb	HalfBathComb	Bedrooms
SalePrice				
count	1460.000000	1460.000000	1460.000000	1460.000000
mean	10516.828082	1.990411	0.440411	2.866438
std	9981.264932	0.732046	0.554016	0.815778
min	1300.000000	0.000000	0.000000	0.000000
25%	7553.500000	1.000000	0.000000	2.000000
50%	9478.500000	2.000000	0.000000	3.000000
75%	11601.500000	2.000000	1.000000	3.000000
max	215245.000000	6.000000	4.000000	8.000000

We can see that there is a huge difference in the mean and deviation values of input parameters (LotArea, FullBathComb, HalfBathComb & Bedrooms). We will apply z-score normalization to normalize these values


```

train_data_transf =
StandardScaler().fit_transform(train_data_filt.iloc[:, :-1])
train_data_transf

array([[ -0.20714171,  1.37960564,  1.01040607,  0.16377912],
       [ -0.09188637,  0.01310345,  1.01040607,  0.16377912],
       [  0.07347998,  1.37960564,  1.01040607,  0.16377912],
       ...,
       [ -0.14781027,  0.01310345, -0.79521555,  1.39002276],
       [ -0.08016039,  0.01310345, -0.79521555, -1.06246453],
       [ -0.05811155,  0.01310345,  1.01040607,  0.16377912]])

x_train_transf = pd.DataFrame(train_data_transf, columns=['LotArea',
'FullBathComb', 'HalfBathComb', 'Bedrooms'])
x_train_transf

   LotArea  FullBathComb  HalfBathComb  Bedrooms
0   -0.207142      1.379606      1.010406    0.163779
1   -0.091886      0.013103      1.010406    0.163779
2    0.073480      1.379606      1.010406    0.163779
3   -0.096897      0.013103     -0.795216    0.163779
4    0.375148      1.379606      1.010406    1.390023
...      ...      ...      ...      ...
1455 -0.260560      0.013103      1.010406    0.163779
1456  0.266407      1.379606     -0.795216    0.163779
1457 -0.147810      0.013103     -0.795216    1.390023
1458 -0.080160      0.013103     -0.795216   -1.062465
1459 -0.058112      0.013103      1.010406    0.163779

[1460 rows x 4 columns]

y_train = train_data_filt['SalePrice']
y_train

0      208500
1      181500
2      223500
3      140000
4      250000
...
1455     175000
1456     210000
1457     266500
1458     142125
1459     147500
Name: SalePrice, Length: 1460, dtype: int64

X_train, X_test, y_train, y_test = train_test_split(
    x_train_transf, y_train, test_size=0.25, random_state=42)

```

We will create the linear regression model using sklearn's LinearRegression class module and fit it to our housing data

```
lr_model = LinearRegression()
lr_model.fit(X_train, y_train)
lr_model
```

```
LinearRegression()
```

#R2 Value

```
lr_model.score(X_test, y_test)
```

```
0.4529185974643213
```

```
test_data
```

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley
LotShape \							
0	1461	20	RH	80.0	11622	Pave	NaN
Reg							
1	1462	20	RL	81.0	14267	Pave	NaN
IR1							
2	1463	60	RL	74.0	13830	Pave	NaN
IR1							
3	1464	60	RL	78.0	9978	Pave	NaN
IR1							
4	1465	120	RL	43.0	5005	Pave	NaN
IR1							
...
...							
1454	2915	160	RM	21.0	1936	Pave	NaN
Reg							
1455	2916	160	RM	21.0	1894	Pave	NaN
Reg							
1456	2917	20	RL	160.0	20000	Pave	NaN
Reg							
1457	2918	85	RL	62.0	10441	Pave	NaN
Reg							
1458	2919	60	RL	74.0	9627	Pave	NaN
Reg							
	LandContour	Utilities	...	ScreenPorch	PoolArea	PoolQC	Fence \
0	Lvl	AllPub	...	120	0	NaN	MnPrv
1	Lvl	AllPub	...	0	0	NaN	NaN
2	Lvl	AllPub	...	0	0	NaN	MnPrv
3	Lvl	AllPub	...	0	0	NaN	NaN
4	HLS	AllPub	...	144	0	NaN	NaN
...
1454	Lvl	AllPub	...	0	0	NaN	NaN
1455	Lvl	AllPub	...	0	0	NaN	NaN
1456	Lvl	AllPub	...	0	0	NaN	NaN

1457	Lvl	AllPub	...	0	0	NaN	MnPrv
1458	Lvl	AllPub	...	0	0	NaN	NaN

	MiscFeature	MiscVal	MoSold	YrSold	SaleType	SaleCondition
0	NaN	0	6	2010	WD	Normal
1	Gar2	12500	6	2010	WD	Normal
2	NaN	0	3	2010	WD	Normal
3	NaN	0	6	2010	WD	Normal
4	NaN	0	1	2010	WD	Normal
...
1454	NaN	0	6	2006	WD	Normal
1455	NaN	0	4	2006	WD	Abnorml
1456	NaN	0	9	2006	WD	Abnorml
1457	Shed	700	7	2006	WD	Normal
1458	NaN	0	11	2006	WD	Normal

[1459 rows x 80 columns]

Testing the model performance

```
test_data = pd.read_csv('test.csv')
fullcomb_test = test_data['FullBath'] + test_data['BsmtFullBath']
halfcomb_test = test_data['HalfBath'] + test_data['BsmtHalfBath']
test_data_x = pd.DataFrame({'LotArea':test_data['LotArea'],
'FullBathComb':fullcomb_test, 'HalfBathComb':halfcomb_test,
'Bedrooms':test_data['BedroomAbvGr']})
test_data_x
```

	LotArea	FullBathComb	HalfBathComb	Bedrooms
0	11622	1.0	0.0	2
1	14267	1.0	1.0	3
2	13830	2.0	1.0	3
3	9978	2.0	1.0	3
4	5005	2.0	0.0	2
...
1454	1936	1.0	1.0	3
1455	1894	1.0	1.0	3
1456	20000	2.0	0.0	4
1457	10441	1.0	1.0	3
1458	9627	2.0	1.0	3

[1459 rows x 4 columns]

```
test_data_transf =
StandardScaler().fit_transform(test_data_x.iloc[:, :])
test_data_transf
```

```
array([[ 0.36392912, -1.29367496, -0.82114784, -1.02954254],
       [ 0.89786065, -1.29367496,  1.03375511,  0.17599724],
       [ 0.80964587, -0.0061856 ,  1.03375511,  0.17599724],
       ...,
```

```
[ 2.05514965, -0.0061856 , -0.82114784,  1.38153702],
[ 0.12552719, -1.29367496,  1.03375511,  0.17599724],
[-0.03879049, -0.0061856 ,  1.03375511,  0.17599724]])
```

```
x_test_transf = pd.DataFrame(test_data_transf, columns=['LotArea',
'FullBathComb', 'HalfBathComb', 'Bedrooms'])
```

```
x_test_transf
```

	LotArea	FullBathComb	HalfBathComb	Bedrooms
0	0.363929	-1.293675	-0.821148	-1.029543
1	0.897861	-1.293675	1.033755	0.175997
2	0.809646	-0.006186	1.033755	0.175997
3	0.032064	-0.006186	1.033755	0.175997
4	-0.971808	-0.006186	-0.821148	-1.029543
...
1454	-1.591330	-1.293675	1.033755	0.175997
1455	-1.599808	-1.293675	1.033755	0.175997
1456	2.055150	-0.006186	-0.821148	1.381537
1457	0.125527	-1.293675	1.033755	0.175997
1458	-0.038790	-0.006186	1.033755	0.175997

```
[1459 rows x 4 columns]
```

```
x_test_transf = x_test_transf.fillna(0)
```

```
y_predict = lr_model.predict(x_test_transf)
```

```
prediction_data = pd.DataFrame({'Id':test_data['Id'],
'SalePrice':y_predict})
```

```
prediction_data
```

	Id	SalePrice
0	1461	114575.795619
1	1462	153043.709580
2	1463	207161.466551
3	1464	199535.514487
4	1465	156458.769683
...
1454	2915	128631.556177
1455	2916	128548.407167
1456	2917	186270.587236
1457	2918	145469.230712
1458	2919	198840.626331

```
[1459 rows x 2 columns]
```

```
prediction_data.to_csv('prediction.csv', index=False)
```